

REVEAL — System Architecture (Plain-English Blueprint)


A visual guide to how REVEAL is built, written for founders, investors, and partners — not just engineers. No prior software knowledge assumed.

1. What REVEAL is, in one breath

REVEAL is a web application that helps design students discover **which kind of designer they actually are** — by reading *what they do*, not just *what they say they like*.

A student signs up, works through a guided survey of scenarios and small tasks, and — after a human on our team approves it — receives a polished, personalised **Design Signature** report: charts, plain-language insight, and tailored growth experiments.

The whole product is three moving parts working together:



Everything else in this document explains those three boxes and how a student's data travels between them.

2. The one big idea (why the architecture matters)

Most "AI report" products let a language model *decide* what's true about you and *also* write it up. That is fast to build — and impossible to trust, because the same answer can come out differently every time.

REVEAL is deliberately built the other way around. We split the job in two:

The job	Who does it	Why it's split out
Deciding what's true about a student	Plain, fixed math (no AI)	Same inputs → <i>identical</i> result, every single time. Auditable. Defensible.
Phrasing it nicely into a report	One AI call, tightly boxed in	The AI only chooses <i>words</i> , never <i>verdicts</i> . It can't change the score.

The investor takeaway: the part of REVEAL that makes a *claim about a person* is deterministic and reproducible. The AI is a writer, not a judge. That is the core of our credibility — and it's baked into the architecture, not bolted on as a policy.

3. The technology, translated

Here is every major piece of technology REVEAL runs on, and what each one *is* in everyday terms.

The website the student sees (`apps/web`)

Technology	What it actually is	Everyday analogy
React	The framework that builds the interactive screens	The set of LEGO bricks the storefront is built from
Vite	The tool that bundles and serves those screens	The factory that snaps the bricks into a finished shop
Tailwind CSS	The styling system (colours, spacing, the "premium" look)	The interior designer who makes every room match
React Router	Moves the student between pages without reloading	The hallways connecting the rooms
React Query / Zustand	Remembers data and app state while you use it	The shop's short-term memory during your visit

The output is a **static website** — plain files (HTML, JavaScript, images) that can be served from anywhere in the world, cheaply and instantly.

The server that does the thinking (`apps/api`)

Technology	What it actually is	Everyday analogy
Node.js	The engine that runs our server-side code	The electricity that powers the back office
Fastify	The framework that answers web requests	The receptionist who takes every request and routes it
The Analysis Engine	Our own pure-math scoring code	The examiner who grades the work by a fixed rubric
Anthropic Claude	The single AI call that writes the prose	The copywriter who turns grades into a readable report
bcrypt + JWT	Password protection and login passes	The lock on the door and the wristband you wear inside

The shared rulebook (`packages/shared`)

Technology	What it actually is	Everyday analogy
Zod contracts	Strict definitions of what every piece of data must look like	The customs officer who rejects any malformed parcel
Frozen enums & constants	The fixed vocabulary and scoring numbers	The official dictionary both sides must speak


Because the website *and* the server both import this one rulebook, they can never drift out of sync. A question added on one side is understood by the other automatically.

Where everything is remembered

Technology	What it actually is	Everyday analogy
PostgreSQL	The database — a reliable, structured filing system	The bank vault with numbered, indexed safety-deposit boxes

4. The four layers (the heart of the design)

Inside the server, a student's data flows through **four sealed stages**. Each stage only ever hands its result forward — it never reaches back. This is what makes the whole thing predictable and safe.



```
RAW ANSWERS → [1 CAPTURE] → [2 ENGINE: the math] →  HUMAN APPROVES
                                     |
                                     ↓
                                [3 AI writes it up] ←
                                     |
                                [4 SAVED ONCE] → shown to student
```


What each layer guarantees

Layer	What goes in	What comes out	The promise it makes
1 · Capture	Taps, choices, timings	Raw records, untouched	We never quietly alter what you told us
2 · Engine	Raw records	A "Findings Object" (the scores)	Same input → same findings, forever
3 · Synthesis	Findings Object	Report text in fixed slots	The AI writes; it does not decide
4 · Cache	Report text	The saved report	Generated once, then only ever re-served

The database enforces the "once" promise physically: the report table can hold **exactly one** report per student. A second attempt to write is impossible, not just discouraged.

5. A student's journey, end to end

This is the complete trip a single student's data takes, from first click to finished report.



Narrated, step by step:

- 1. **Arrive & register.** The student lands on the public site and signs up. Their password is never stored as text — it's scrambled with *bcrypt* so even we can't read it. Email is verified before they can begin.
- 2. **Begin the survey.** The survey is split into short sittings. The student can stop and resume; the server always knows the exact next question (its "cursor").
- 3. **Answer by answer.** Every choice, and even how long it took, is sent to the server and filed in the raw-capture vault — **exactly as given, never edited.**
- 4. **Seal the session.** When a sitting is complete the student seals it. Sealed answers are locked by the database itself and can't be altered afterward.

- 5. **The math runs.** Once capture is complete, the Analysis Engine turns the raw answers into a Findings Object — the student's capacities, values, working style, and the conditions their work needs. **No AI is involved here.**
- 6. **It waits for a human.** Crucially, the report is *not* generated yet. The instance sits in a queue for our team.
- 7. **A team member approves.** An admin reviews the findings — especially any "high-stakes" or surprising results — and approves.
- 8. **The AI writes, once.** Only now does a single Claude call run, turning the findings into polished prose that fills a fixed set of report slots. The result is saved permanently.
- 9. **The student reads it.** From this point on, opening the report just re-serves the saved copy. It never regenerates, so it never changes underneath them.

If we ever have no AI key configured (or set `SYNTHESIS_MODE>manual`), the system falls back to a built-in deterministic writer. **The entire product works end-to-end with zero AI calls** — the AI is an enhancement, not a dependency.

6. What we actually store (the data model)

The database is organised as a clean "spine" — one student has one report instance, which has one raw-capture record, which produces one findings record, which produces one report. Everything hangs off that line.




Table	Holds	Layer
student	Account, profile, consent	—
report_instance	One attempt; tracks status	spine

Table	Holds	Layer
raw_capture	Every raw answer, untouched	1 · Capture
derived	The Findings Object (the scores)	2 · Engine
report_payload	The final report — one row max	4 · Cache
review	Who approved, and their notes	Human gate
instrument_*	The questions/tasks themselves, versioned & admin-editable	Content

Two design details worth calling out to a technical partner:

- **The survey content lives in the database, versioned.** Our team can edit questions through an admin portal without a developer or a redeploy.
- **Sealed answers are protected by the database itself** — a built-in rule rejects any attempt to change an answer once its session is sealed. Integrity isn't left to trust; it's enforced at the vault door.

7. Where the AI is — and, deliberately, isn't



Every judgement REVEAL makes about a person is produced by the engine and can be recomputed and audited at any time. The AI never sees a student's raw data as a blank canvas — it receives a finished, structured verdict and is asked only to express it well. This is the single most important architectural decision in the product, and it is what lets us stand behind every report.

8. Security & privacy posture (at a glance)

Concern	How the architecture handles it
Passwords	Never stored as text — hashed with bcrypt, one-way
Login sessions	Signed JWT passes with an expiry; no server-side session to leak
Who can do what	Students and admins have separate, enforced roles (RBAC)
Bad/malicious input	Every request is validated against the shared Zod contracts before it touches the database
Consent	Recorded explicitly per student before capture begins
Data integrity	Sealed answers are immutable; reports generate exactly once
Transport	Database connections use SSL in production

9. The monorepo map (for the technical reader)

Everything lives in one repository, split into three packages managed by `pnpm`.

```
reveal/
├─ packages/shared/  The shared rulebook – enums, scoring constants,
|                    Zod contracts, report slots, scenario definitions.
|                    Imported by BOTH web and api so they never disagree.
|
├─ apps/api/         The Node.js + Fastify server. Contains the four-layer
|   └─ src/engine/    Layer 2 · the pure-math analysis engine
|   └─ src/synthesis/ Layer 3 · the single AI call (+ deterministic fallback)
|   └─ src/routes/    the web endpoints (auth, capture, report, admin...)
|   └─ src/repositories/ the code that reads/writes the database
|   └─ src/db/        schema migrations + seed data
|
├─ apps/web/         The React + Vite + Tailwind website the student uses.
|   └─ src/features/  landing, auth, survey capture, dashboard, report, admin
|
└─ infra/            Ready-made deploy configs (see the deployment plan)
```

The four-layer pipeline described above lives inside `apps/api/src`. The engine is the crown jewel: pure functions with no database or network access, which is exactly why it's so easy to test and trust.

10. Mini-glossary

Term	Plain meaning
Frontend / web	The part you see and click in the browser
Backend / API / server	The part that runs privately and does the work
Database	The permanent, structured store of everything
Deterministic	Same input always gives the same output
Findings Object	The engine's structured verdict about a student
Synthesis	Turning that verdict into readable report prose
Cache	A saved copy that's re-served instead of re-made
Migration	A versioned change to the database's shape
Seed	Loading the starting content (questions, admins)
Monorepo	One codebase holding several related projects

This document describes how REVEAL is built. For how to put it live on the internet at no cost, see [ZERO_COST_DEPLOYMENT_PLAN.md](#).